

Literature review and Implementation of Style Transfer


UIC 统计系

李杰

1730005024

What's style transfer?

- Computer version-Image transfer-**Style-transfer**



Literature review of Style-Transfer

A method of object recognition:

“Bag of feature”




The object we want to recognition



Local feature
(select by human)

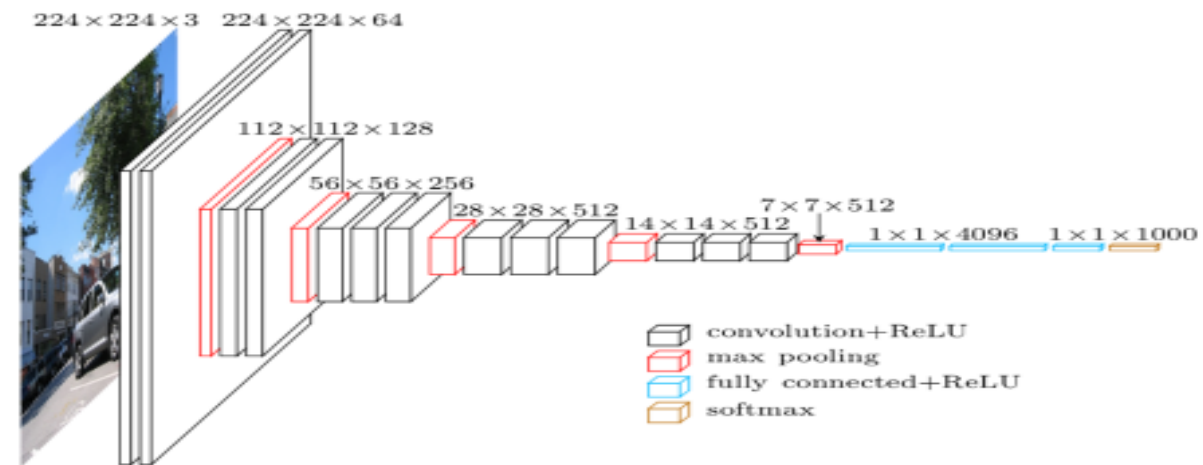
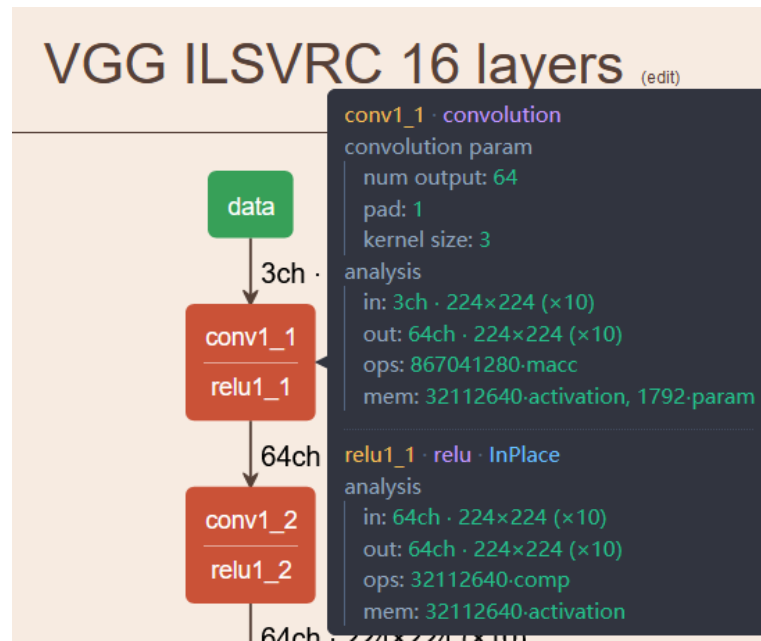
Literature review of Style-Transfer

How to do the object recognition better?



AlexNet(2012)

VGG16(2014)



Literature review of Style-Transfer

- How to use small kernel(VGG16) to replace large kernel(Alex-net) ?





Figure 1. Mini-network replacing the 5×5 convolutions.

Literature review of Style-Transfer

- 2015: The first model

A Neural Algorithm of Artistic Style

Leon A. Gatys,^{1,2,3*} Alexander S. Ecker,^{1,2,4,5} Matthias Bethge^{1,2,4}



Use VGG16 to extract the feature of **Content** and **Style** from two pictures

Literature review of Style-Transfer

• Other important model



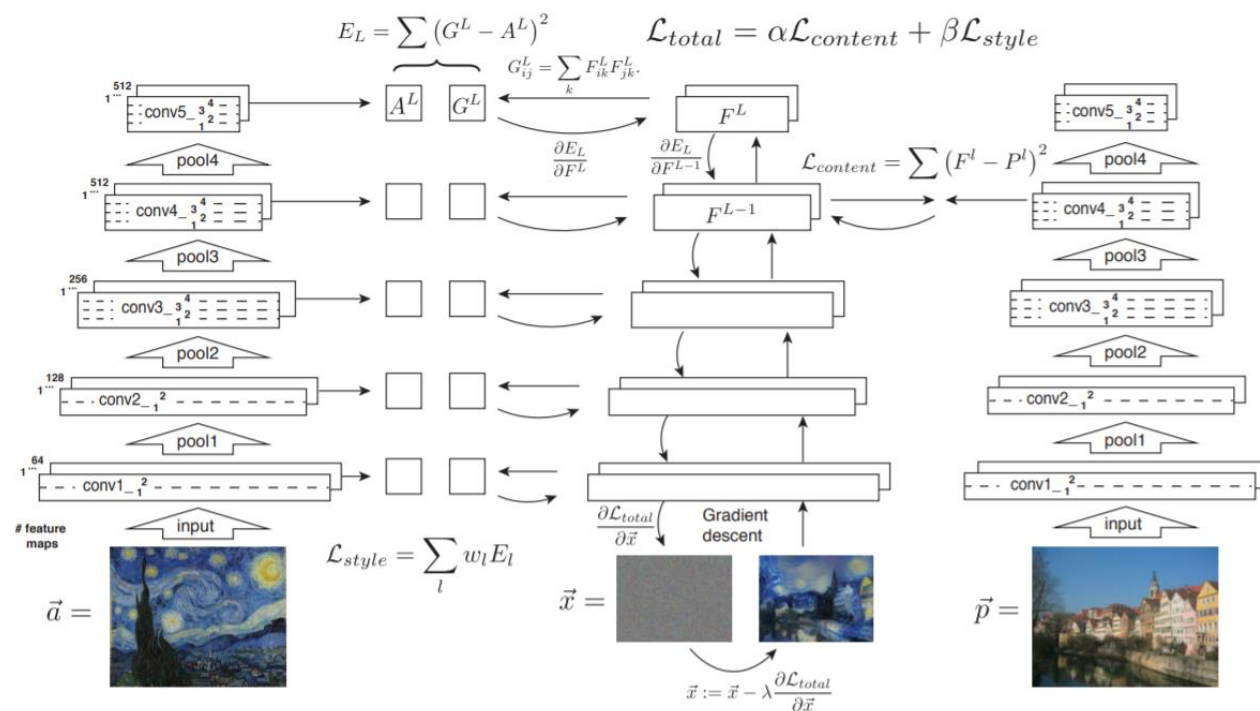

Gatys's model

Two assumption:

Image = Content+Style

Style = Texture

$$\mathcal{L}_{total} = \alpha \mathcal{L}_{content} + \beta \mathcal{L}_{style}$$



Gatys's model

Measure the content:

$$\mathcal{L}_{\text{content}}(\vec{p}, \vec{x}, l) = \frac{1}{2} \sum_{i,j} (F_{ij}^l - P_{ij}^l)^2 . \quad (1)$$

Represent the respond of content in Layer l ($l=1 \sim Nl$), Filter i ($i = 1 \sim Ml$) , Position j

The derivative of this loss with respect to the activations in layer l equals

$$\frac{\partial \mathcal{L}_{\text{content}}}{\partial F_{ij}^l} = \begin{cases} (F^l - P^l)_{ij} & \text{if } F_{ij}^l > 0 \\ 0 & \text{if } F_{ij}^l < 0 , \end{cases} \quad (2)$$

Gatys's model

Measure the style:

suppose that the texture of image can represent the style

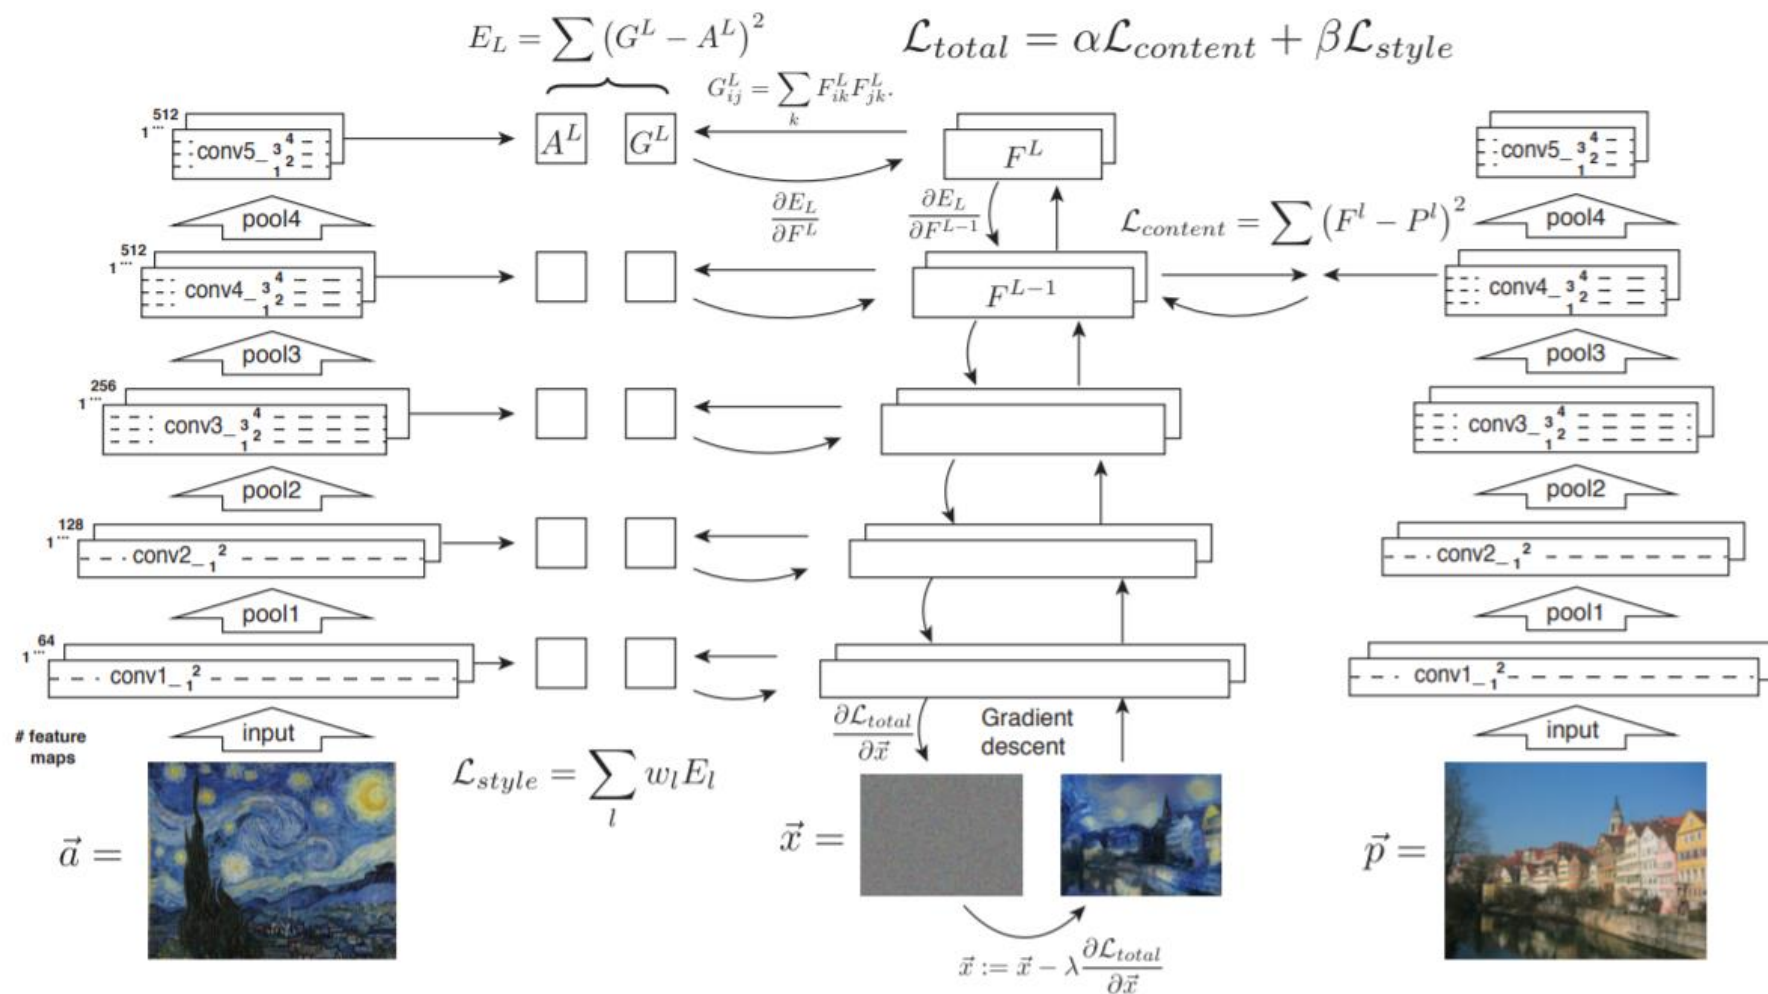
Gram matrix: used to capture the texture information of Image $G_{ij}^l = \sum_k F_{ik}^l F_{jk}^l$.

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{ij}^l - A_{ij}^l)^2$$

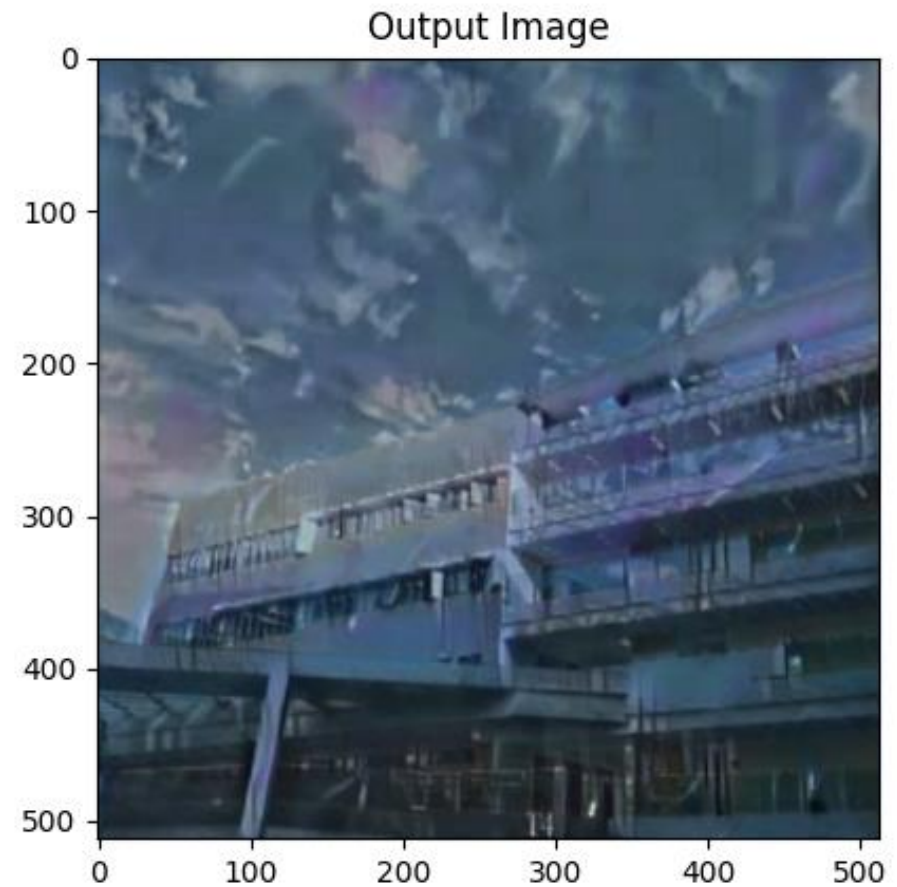
$$\mathcal{L}_{\text{style}}(\vec{a}, \vec{x}) = \sum_{l=0}^L w_l E_l,$$

$$\frac{\partial E_l}{\partial F_{ij}^l} = \begin{cases} \frac{1}{N_l^2 M_l^2} ((F^l)^T (G^l - A^l))_{ji} & \text{if } F_{ij}^l > 0 \\ 0 & \text{if } F_{ij}^l < 0 \end{cases}$$

Gatys's model



Implementation of Gaty's Model (Base on Pytorch)



Style Transfer base on Cycle GANs



Generate Model

Given training data, generate new samples from same distribution



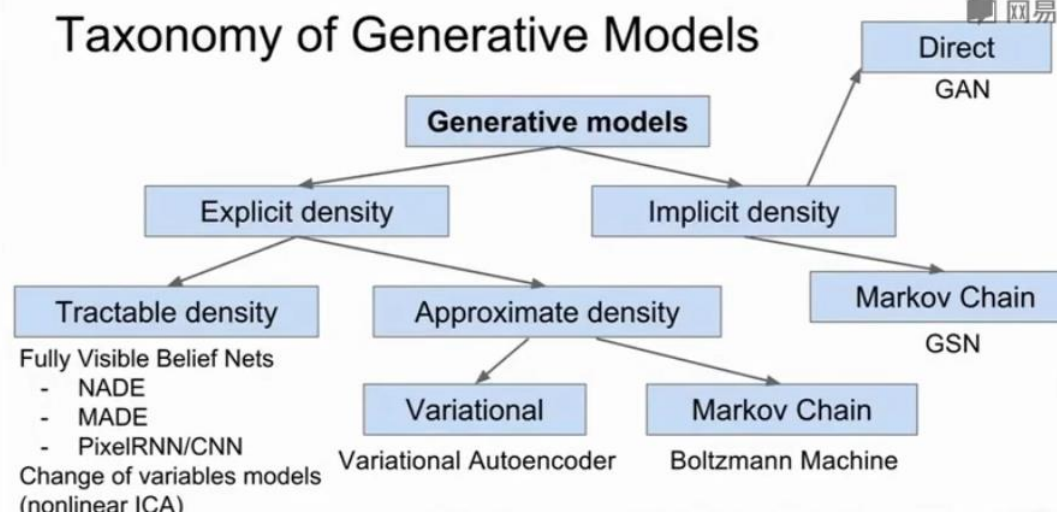
Training data $\sim p_{\text{data}}(x)$



Generated samples $\sim p_{\text{model}}(x)$


Want to learn $p_{\text{model}}(x)$ similar to $p_{\text{data}}(x)$

Taxonomy of Generative Models

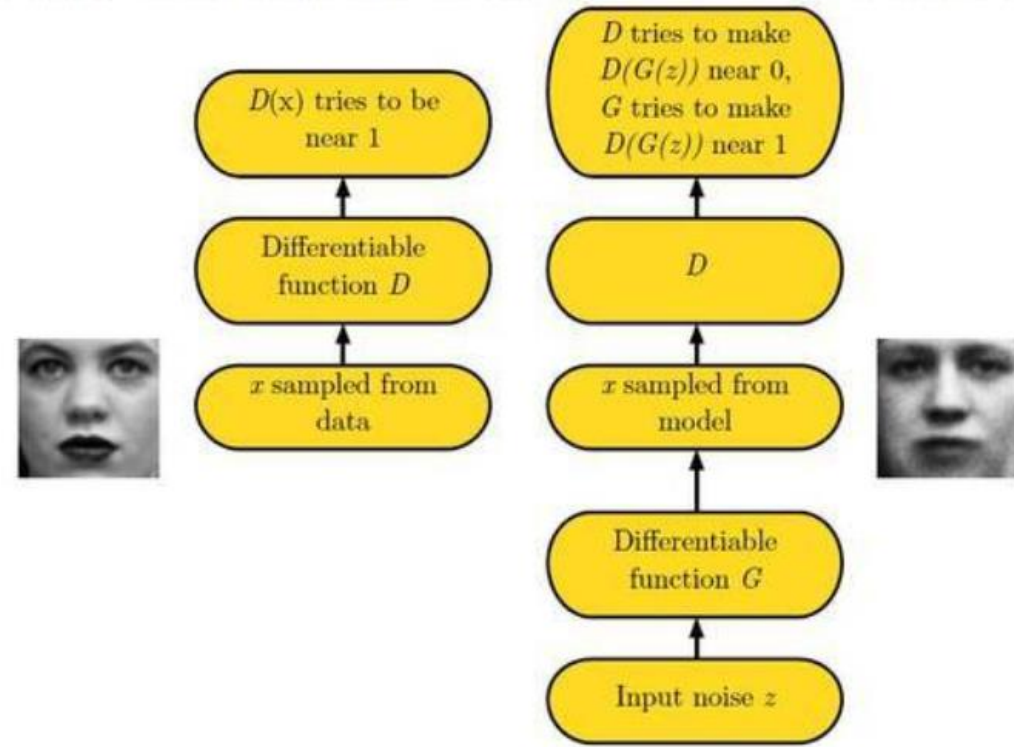


GANs (Ian J. Goodfellow 2014)

(对抗生成网络)



Adversarial Nets Framework



GANs (Ian J. Goodfellow 2014)

(对抗生成网络)

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by **ascending** its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m [\log D(\mathbf{x}^{(i)}) + \log(1 - D(G(\mathbf{z}^{(i)})))]$$

end for

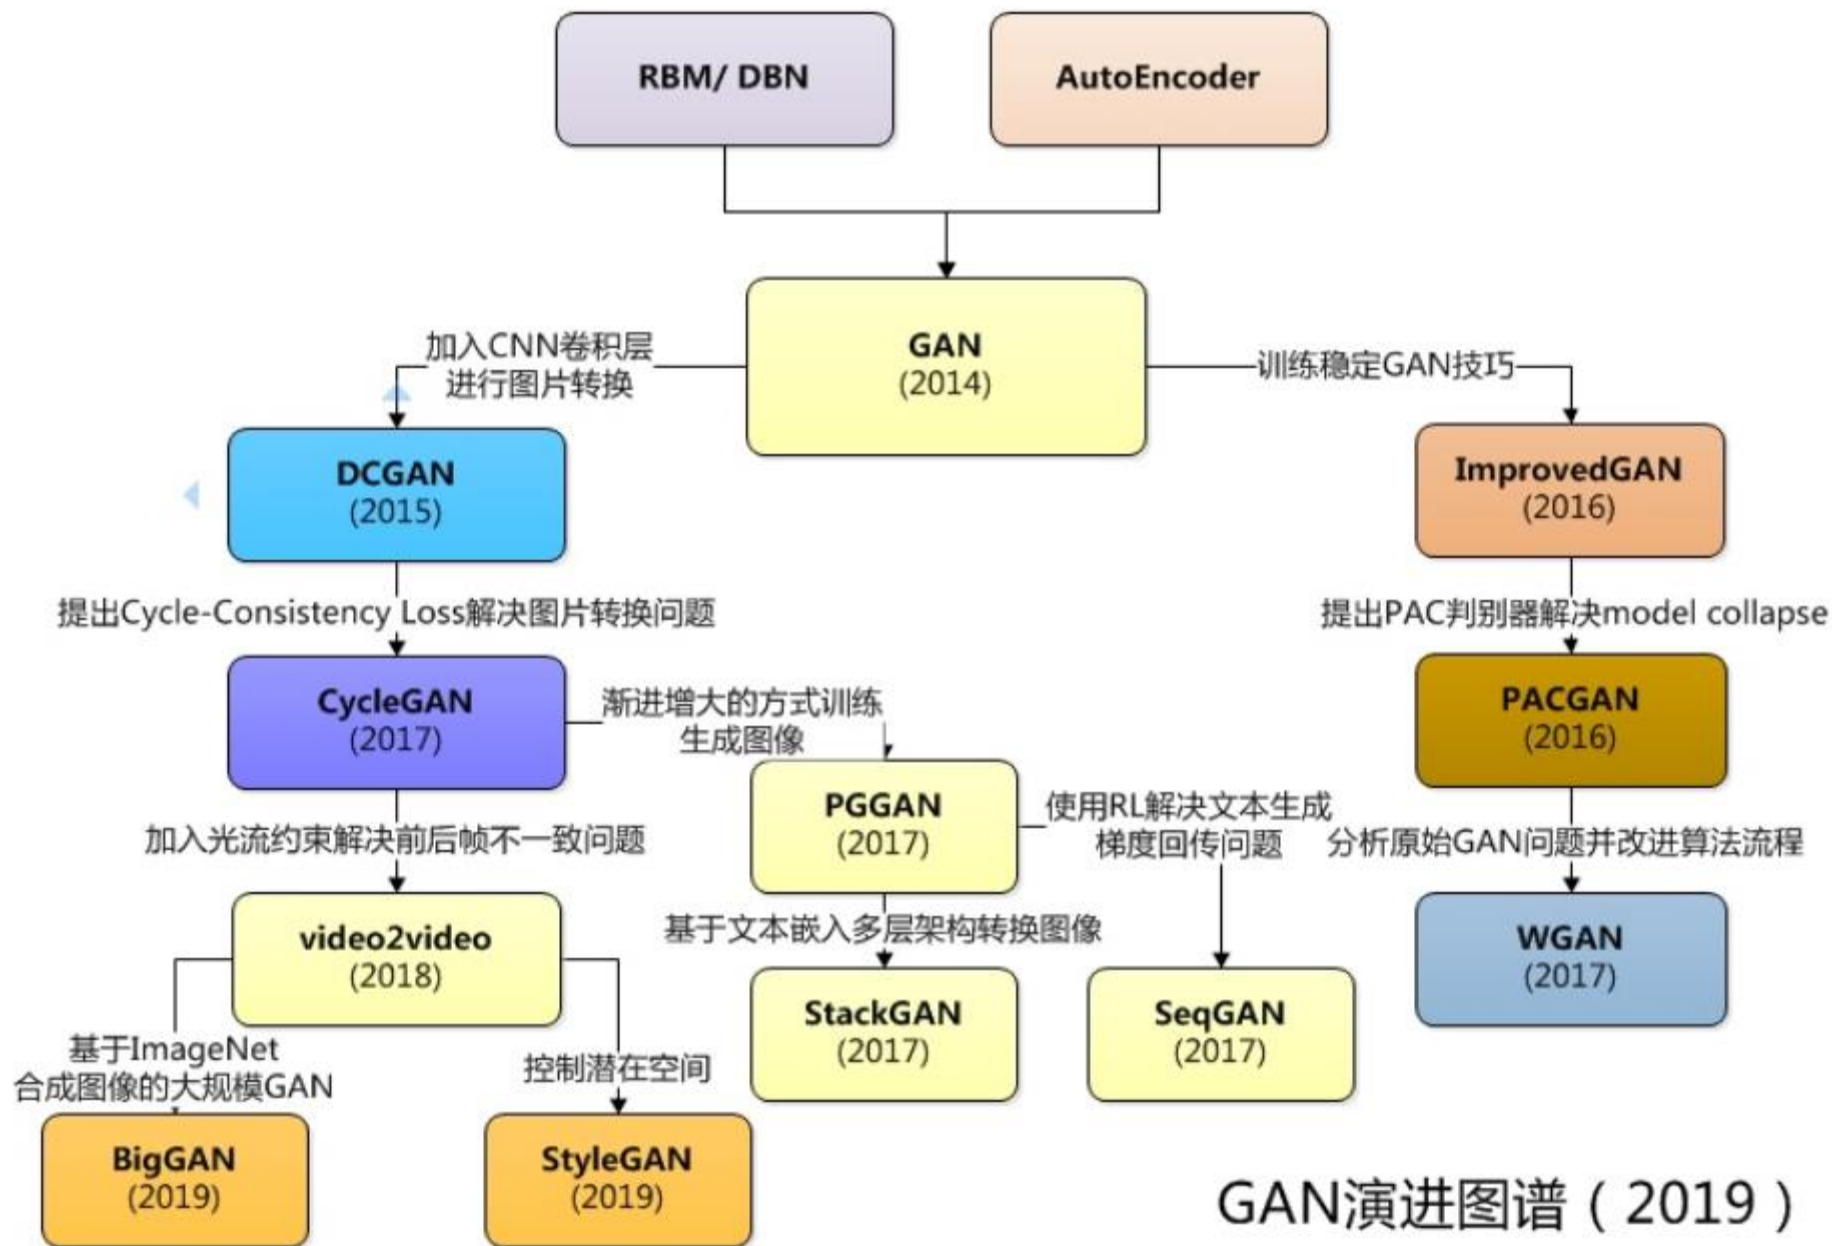
- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by **descending** its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(1 - D(G(\mathbf{z}^{(i)})))$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

$$\nabla C \approx \frac{1}{m} \sum_{j=1}^m \nabla C_{X_j},$$



GAN演进图谱 (2019)

Alec Radford & Luke Metz
Indico Research
Boston, MA
{alec,luke}@indico.io


Soumith Chintala
Facebook AI Research
New York, NY
soumith@fb.com

DCGAN(Deep Convolutional GAN):

Use two CNN as D&G

Architecture guidelines for stable Deep Convolutional GANs



- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.



Why cycle-GANs?

Reason 1: Input for GANs is white noise

Recall the **Purpose of Style transfer** and **GANs**



Why cycle-GANs?

Reason 2: Lack of pair data, the output will be not contain the same content

Image-to-Image Translation with **pix2pix**

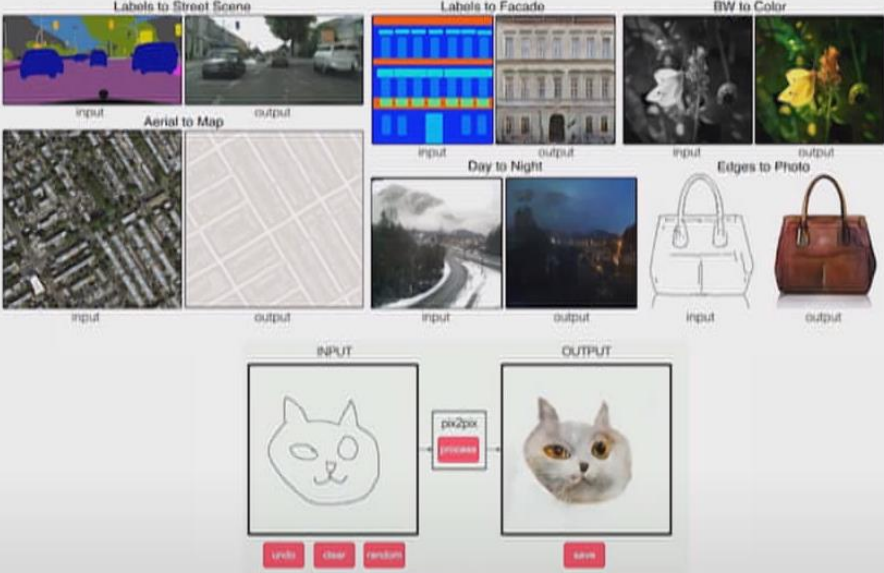
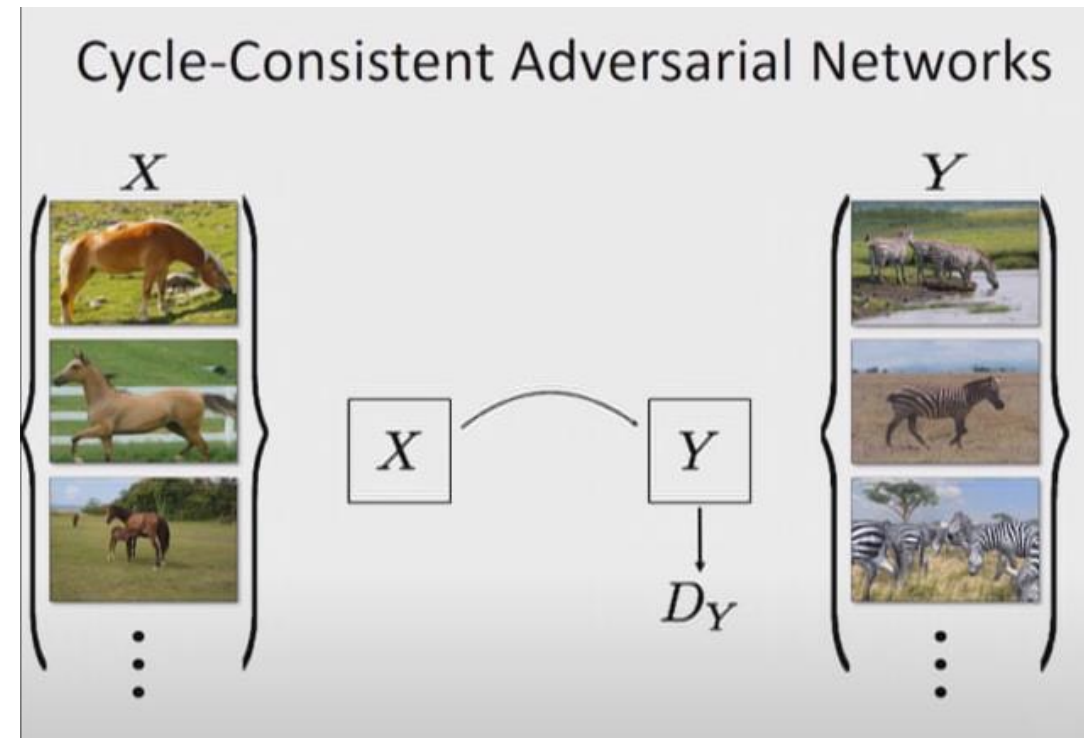





Image-to-image Translation with Conditional Adversarial Nets
Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, Alexei A. Efros. CVPR 2017



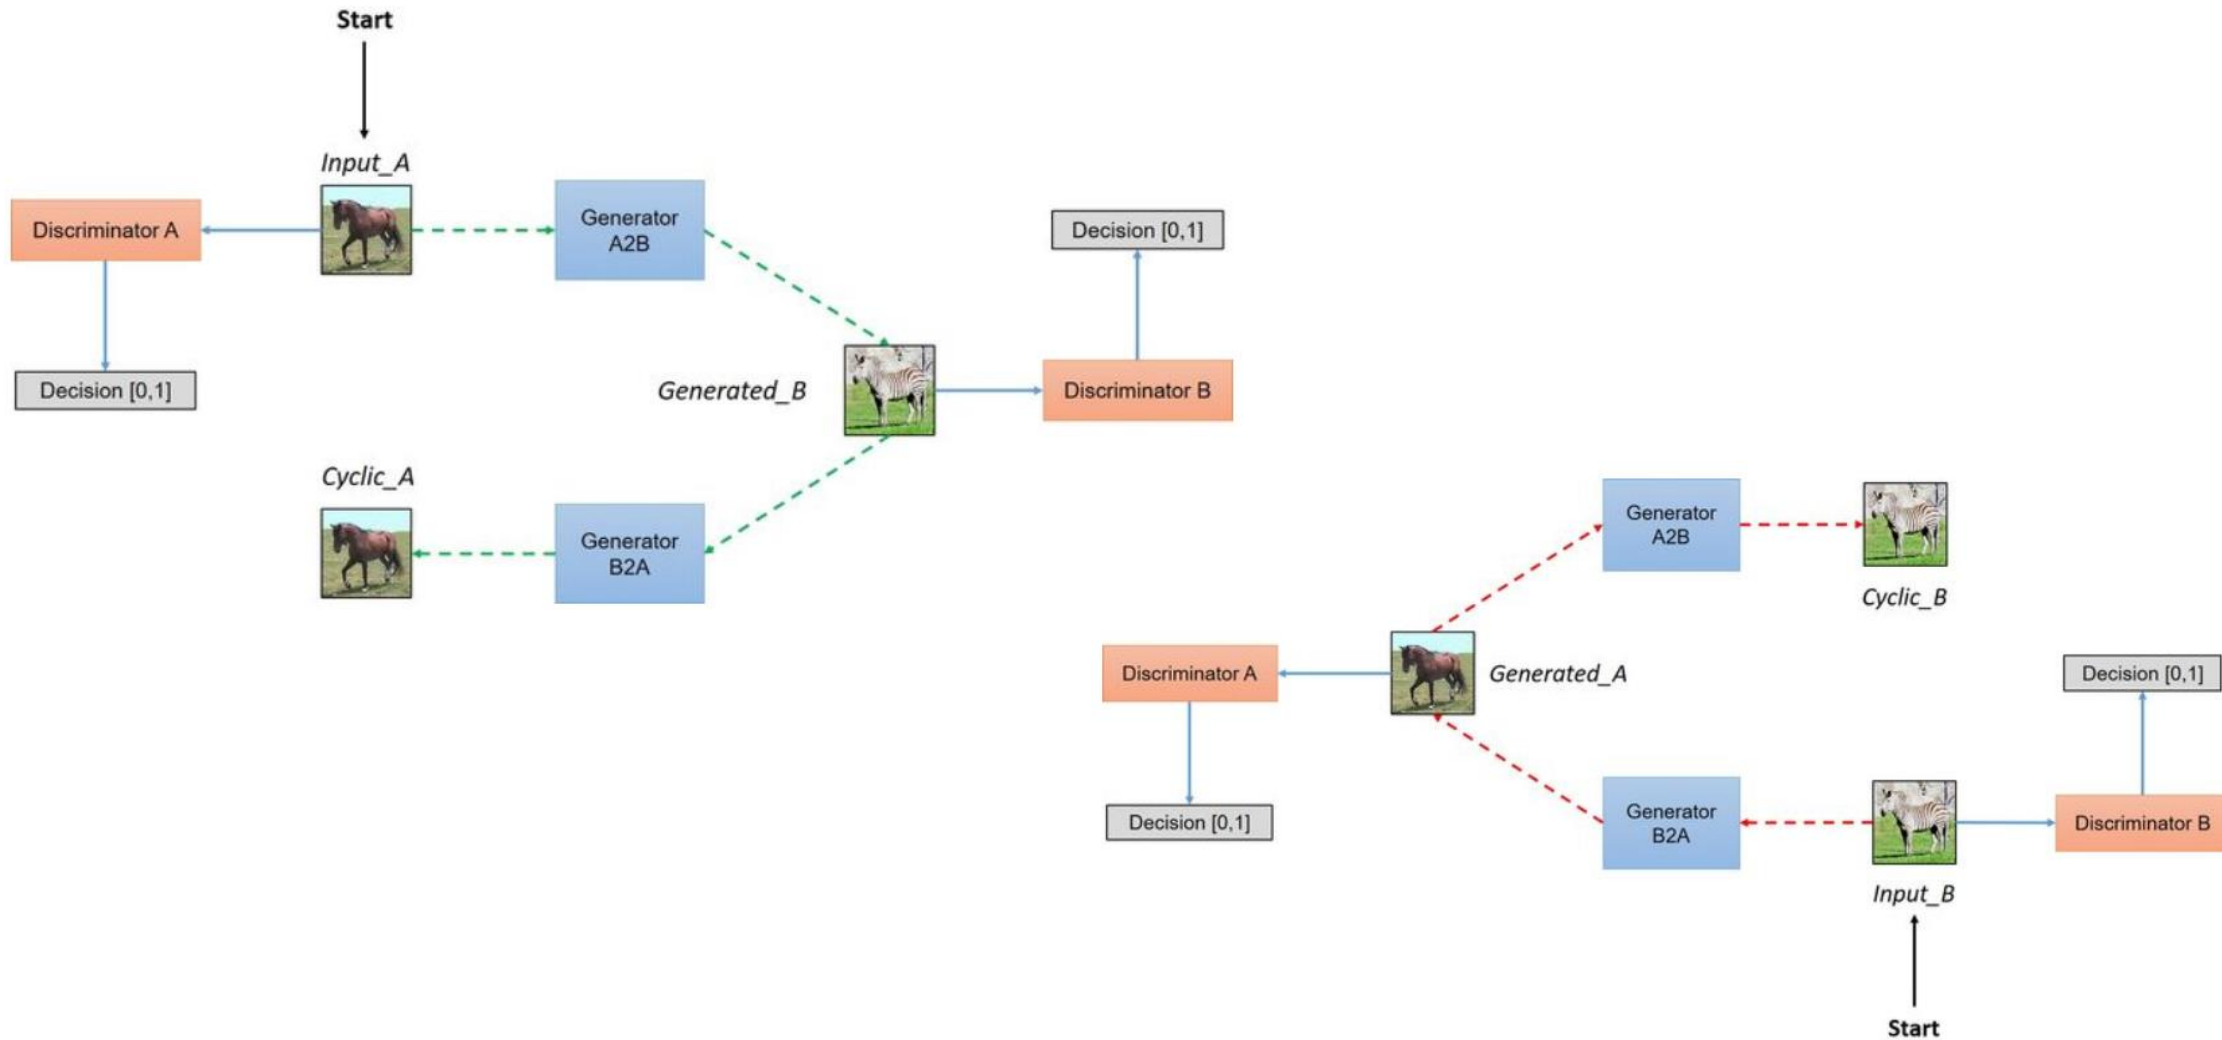
Cycle GAN

Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks

Jun-Yan Zhu* Taesung Park* Phillip Isola Alexei A. Efros
Berkeley AI Research (BAIR) laboratory, UC Berkeley



Cycle GAN



Cycle GAN

$$L_{\text{GAN}}(F, D_Y, X, Y) = E_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] + E_{x \sim p_{\text{data}}(x)} [\log(1 - D_Y(F(x)))]$$

Original Lost function of GANs

$$\begin{aligned} L_{\text{cyc}}(G_{A2B}, G_{B2A}, A, B) &= \mathbb{E}_{x \sim A} [\|G_{B2A}(G_{A2B}(x)) - x\|_1] \\ &+ \mathbb{E}_{y \sim B} [\|G_{A2B}(G_{B2A}(y)) - y\|_1] \end{aligned}$$


$$\begin{aligned} G_{B2A}(G_{A2B}(x)) &\simeq x \\ G_{A2B}(G_{B2A}(y)) &\simeq y \end{aligned}$$

Consider “Consistency Loss”

$$\begin{aligned} L(G_{A2B}, G_{B2A}, D_A, D_B) &= L_G(G_{A2B}, D_B, A, B) \\ &+ L_G(G_{B2A}, D_A, B, A) \\ &+ \lambda L_{\text{cyc}}(G_{A2B}, G_{B2A}, A, B) \end{aligned}$$

Lost function of Cycle-GANs

Result from the paper



Thank you!

- More reference in :

 PyTorch

知乎

arXiv.org